



# Classifying Disaster Tweets

# The problem

## Company

disaster relief  
organizations / news  
agencies

## Context

- Twitter has become an important communication channel in times of emergency.
- The ubiquitousness of smartphones enables people to announce an emergency they're observing in real-time.
- agencies are interested in programatically monitoring Twitter.

## Problem statement

- It's not always clear whether a person's words are actually announcing a disaster.

# Data Wrangling and Exploratory Data Analysis (EDA)



- The data set contains Train and Test sets and taken from Kaggle

<https://www.kaggle.com/c/nlp-getting-started/data>

- Train set has 7613 instances, 4 features, binary target variable.


	id	keyword	location	text	target
7091	10156	upheaval	Connecticut	A look at state actions a year after Ferguson's upheaval <a href="http://t.co/vXUfVT9AU">http://t.co/vXUfVT9AU</a>	1
3002	4313	dust%20storm	NaN	So.... I just watched the trailed for The Dust Storm and I think part of me just died.... Colin is so perfect my goodness.	0
6437	9210	suicide%20bombing	USA	Turkish troops killed in Kurdish militant 'suicide attack' <a href="http://t.co/wD7s6S0vci">http://t.co/wD7s6S0vci</a>	1

- Test set has 3265, same features but no target variable.
- Missing values in both sets for (location and keyword features)




# Inspecting Train Data

- Target Class Distribution



- There are 110 text duplicates and 18 has labeling inconsistency
- Duplicates dropped and inconsistent labels fixed manually

	id	keyword	location	text	target
5641	8044	refugees	NaN	wowo--=== 12000 Nigerian refugees repatriated from Cameroon	0
5620	8018	refugees	NaN	wowo--=== 12000 Nigerian refugees repatriated from Cameroon	1



- Target Class Distribution after cleaning duplicates.

# Inspecting the "location" feature

- 33% of the location feature is missing.
- There are 3341 unique locations for the tweets.
- The naming of location in this data set is not consistent. We see some tweets' location as a country and others as a city.
- the USA is the most frequent location for the tweets, followed by New York, then United States representing (1.37%, 0.93%, and 0.66%) of the tweets respectively.
- **With this large number of missing values, it is not feasible to clean the location values. The location in this data set is not good to be a predictive feature and will not be considered in the modeling.**

	missing values count	%
location	2533	33.272035

	location	count	% of dataset
0	USA	104	1.37
1	New York	71	0.93
2	United States	50	0.66
3	London	45	0.59
4	Canada	29	0.38
...	...	...	...
3336	Republica Dominicana	1	0.01
3337	Republic of the Philippines	1	0.01
3338	Regalo Island	1	0.01
3339	Redondo Beach, CA	1	0.01
3340		1	0.01



# Inspecting the "keyword" feature

- There are 61 (0.8%) missing keywords
- keywords derailment, wreckage, and debris are only appeared in the disaster tweets with disaster probability = 1. Also, outbreak, typhoon, oil%20omb, suicide%20bombing and rescuers has more than 90%
- keywords aftershock, body%20bags, ruin, blazing, body%20bag, electrocute, screaming, traumatised, and blew%20up are mostly appeared in the non disaster tweets, near zero probability of disaster tweets.
- many other keywords that are existed in both classes. The keyword without the context can not be a good predictive of the disaster tweets.

Top 15 keywords with highest probability of indicating disaster tweet

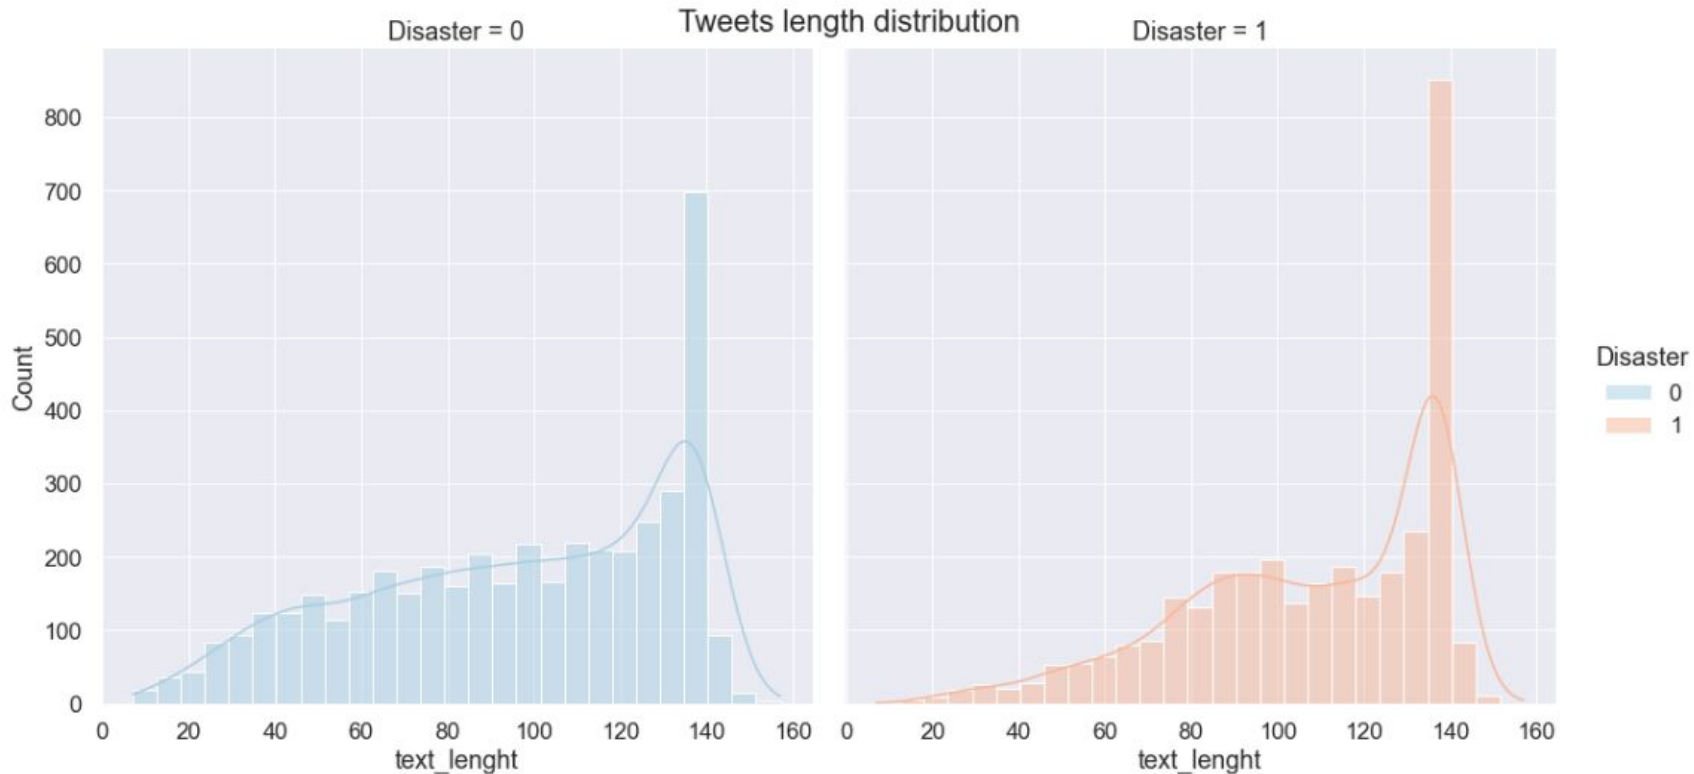
	count	% of dataset rows	Probability of Disaster
keyword			
wreckage	39	0.51	1.000000
debris	37	0.49	1.000000
derailment	39	0.51	1.000000
outbreak	40	0.53	0.975000
oil%20spill	38	0.50	0.973684
typhoon	38	0.50	0.973684
suicide%20bombing	33	0.43	0.969697
suicide%20bomber	31	0.41	0.967742
bombing	29	0.38	0.931034
rescuers	35	0.46	0.914286
suicide%20bomb	35	0.46	0.914286
nuclear%20disaster	34	0.45	0.911765
evacuated	36	0.47	0.888889
razed	35	0.46	0.885714
wildfire	33	0.43	0.878788

The Percentage Count of "keyword" in training set



# Tweets Metadata Analysis

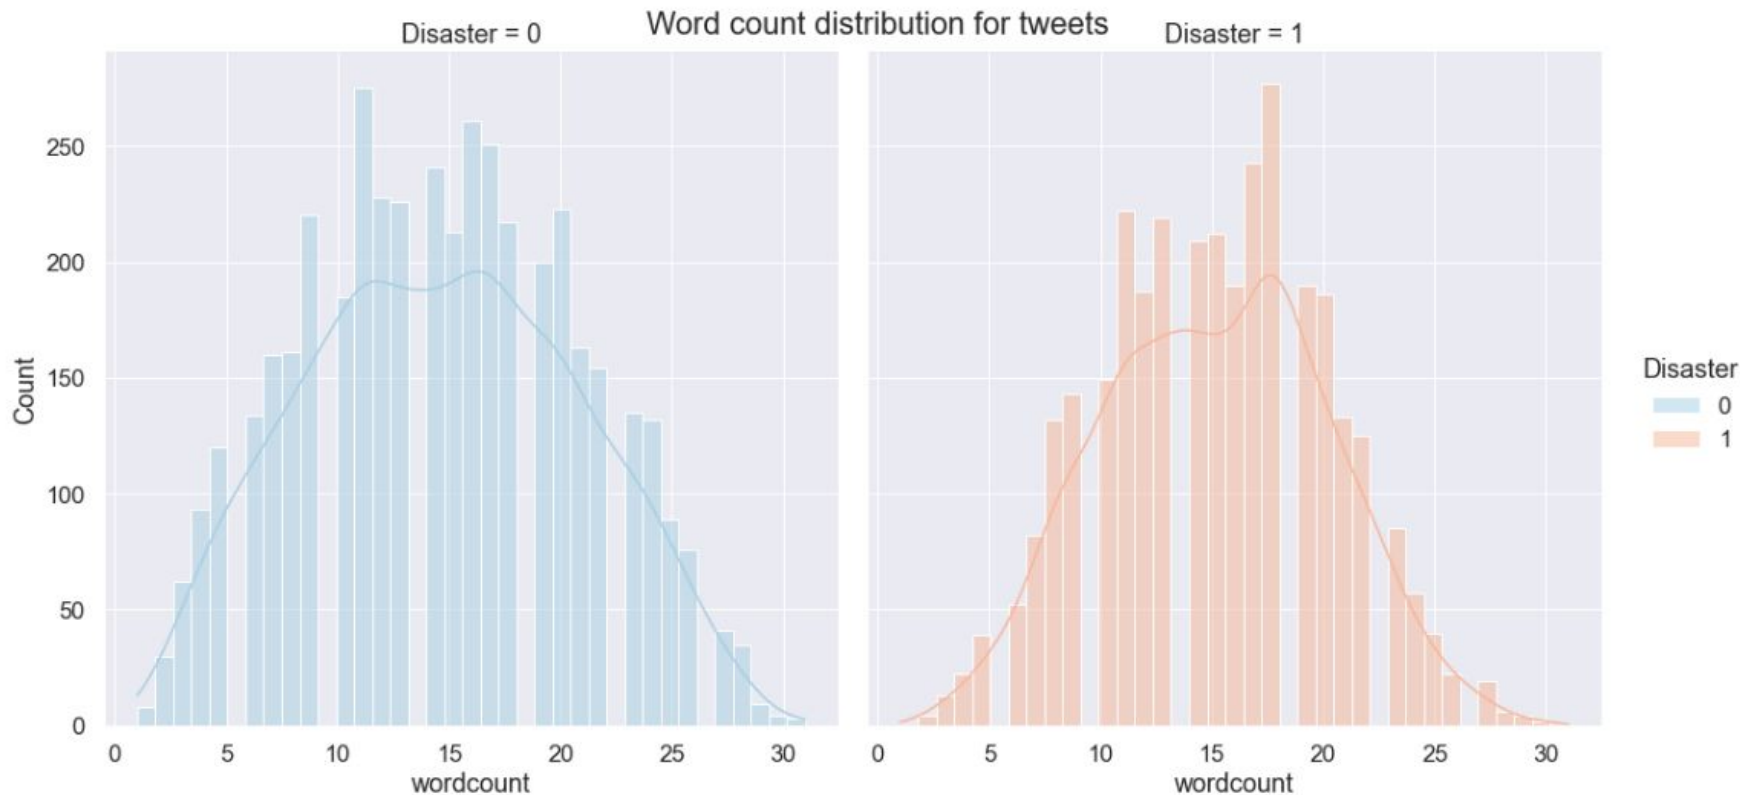
- Tweets length (number of characters) distribution for Disaster and None Disaster classes





# Tweets Meta Data Analysis

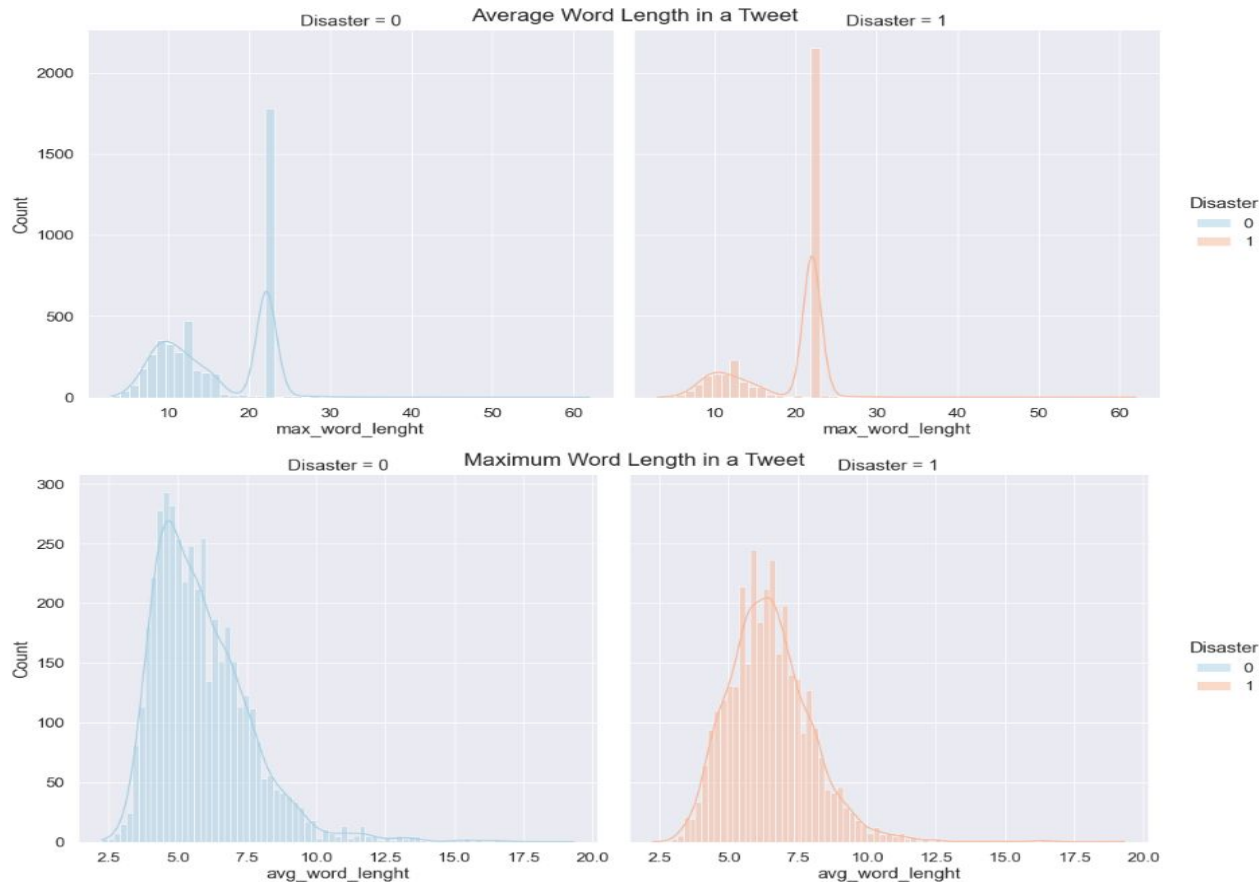
- Word count distribution for tweets



# Tweets Metadata Analysis

- Average and Maximum word length in a tweet


- The Meta data features of both classes have a similar distribution
- None of these Metadata features seems to be a good predictor of the disaster tweets
- We will further check the predictive power of these features



# Exploring Common Words and Phrases

- The distribution of top bigrams after removing stop words

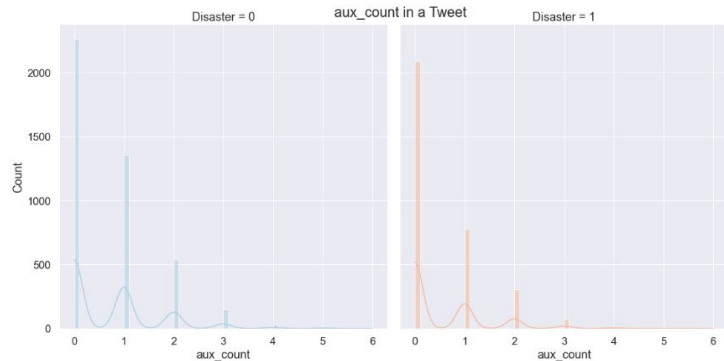
- The data needs to be cleaned from URLs and stop words should be removed to have a meaningful distribution of common phrases in both classes
- There is a noticeable difference of N-Gram output for each class



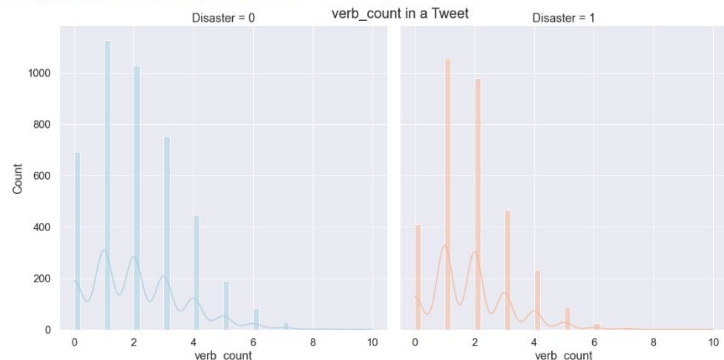
# Exploring Linguistic Features

- Part of Speech Analysis using Spacy

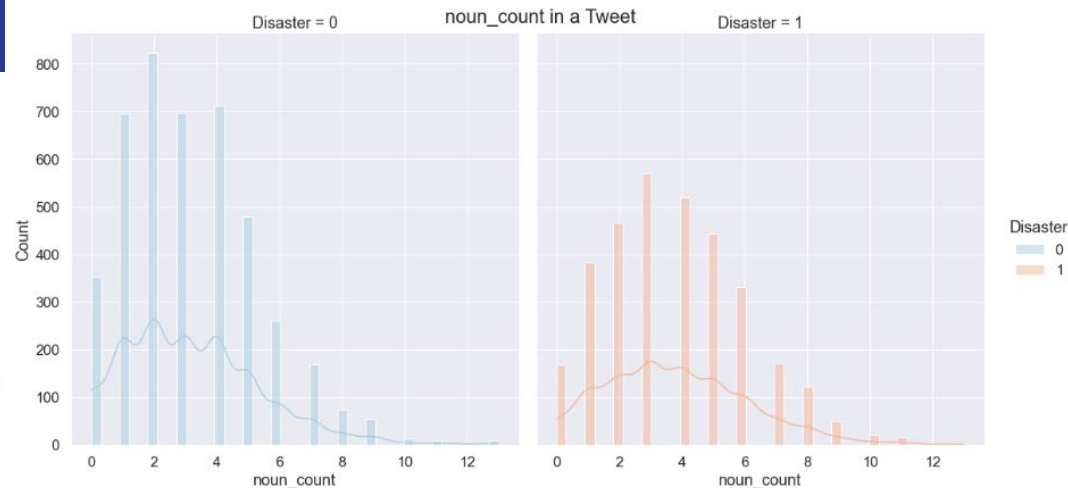
Average aux\_count for disaster tweet is 0.5186887254901961  
Average aux\_count for none disaster tweet is 0.7052195907105082



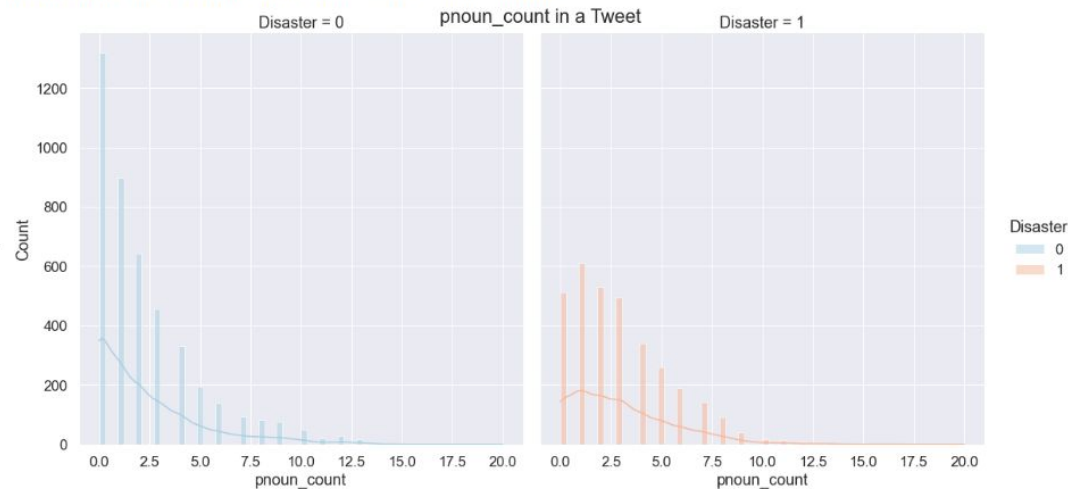
Average verb\_count for disaster tweet is 1.8385416666666667  
Average verb\_count for none disaster tweet is 2.0531156587721315



Average noun\_count for disaster tweet is 3.792892156862745  
Average noun\_count for none disaster tweet is 3.191998160496666



Average pnoun\_count for disaster tweet is 2.9607843137254903  
Average pnoun\_count for none disaster tweet is 2.300988733042079



# 1. Predictive Power of the Imputed Features

- As we noticed from the distribution analysis of the metadata features, the PPS confirms that the computed meta features have no predictive power to the target variable (PPS = 0 for all variables), "disaster" in this case.
- Pearson correlation showed some weak correlation of these metadata variables with the maximum word length being the highest ( $r = 0.25$ ).
- However, Pearson correlation can give misleading results for a binary classification problem (i.e. examining binary target variable), at least, it is not good option. [Click this for more details.](#)

Based on PPS values, these variables should not be used in modeling.





# Data Pre-Processing and Modeling

# Modeling Approaches

## Bag of Words (TF-IDF)

- Multinomial Naive Bayes
- Logistic Regression
- XGboost with Bayesian Hyperparameter Optimization
- Deep Learning:
  - Simple Neural Network
  - Recurrent Neural Network (LSTM)
  - Convolutional Neural Network (CNN)

## Transfer Learning / Deep Learning

- Bidirectional Encoder Representations from Transformers (BERT):
  - Distilled BERT uncased
  - BERT base uncased

# Term Frequency – Inverse Document Frequency

## Pre-processing

- Remove stop word
- Remove urls
- Remove digits
- Remove hashtags
- Remove tags
- Lower case text
- Text lemmatization

## Multinomial Naive Bayes


Accuracy: 0.8

Auc: 0.86

Detail:

	precision	recall	f1-score	support
0	0.77	0.93	0.84	1305
1	0.88	0.63	0.73	979
accuracy			0.80	2284
macro avg	0.82	0.78	0.79	2284
weighted avg	0.82	0.80	0.80	2284

Confusion matrix





# Term Frequency – Inverse Document Frequency

## Logistic Regression


Accuracy: 0.81

Auc: 0.86

Detail:

	precision	recall	f1-score	support
0	0.80	0.90	0.84	1305
1	0.84	0.69	0.76	979
accuracy			0.81	2284
macro avg	0.82	0.80	0.80	2284
weighted avg	0.81	0.81	0.81	2284

Confusion matrix



## XGboost / Bayesian Hyperparameter Optimization


Accuracy: 0.74

Auc: 0.82

Detail:

	precision	recall	f1-score	support
0	0.81	0.71	0.76	1305
1	0.67	0.77	0.72	979
accuracy			0.74	2284
macro avg	0.74	0.74	0.74	2284
weighted avg	0.75	0.74	0.74	2284

Confusion matrix

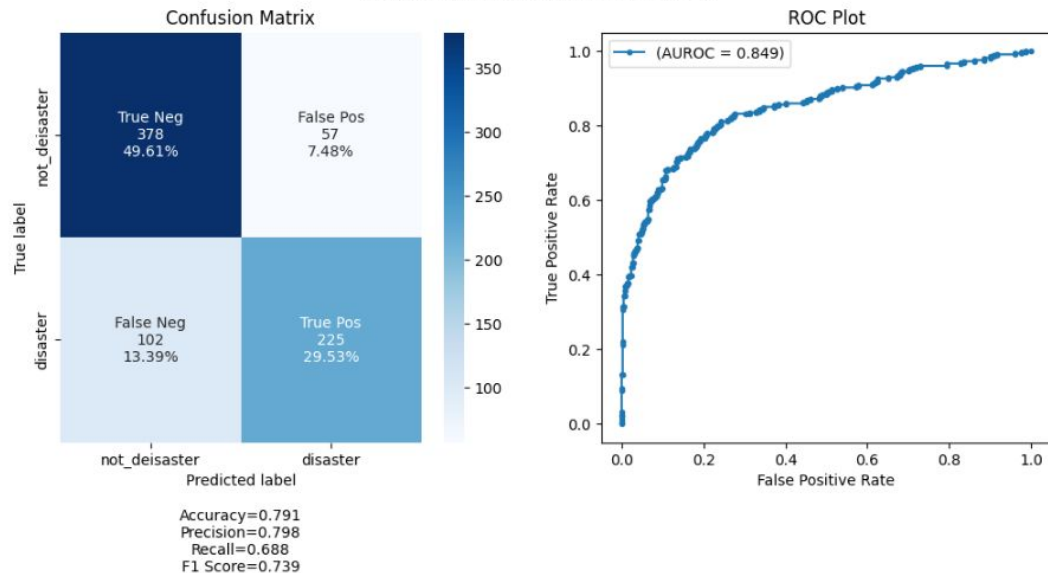


# Term Frequency – Inverse Document Frequency

## Recurrent Neural Network with LSTM with Dropout

## Model Architecture

Model Performance on the Test Dataset



Details:

	precision	recall	f1-score	support
0	0.79	0.87	0.83	435
1	0.80	0.69	0.74	327
accuracy			0.79	762
macro avg	0.79	0.78	0.78	762
weighted avg	0.79	0.79	0.79	762


Model: "sequential\_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 23, 32)	472768
lstm (LSTM)	(None, 60)	22320
dense_4 (Dense)	(None, 1)	61
Total params: 495,149		
Trainable params: 495,149		
Non-trainable params: 0		

# Term Frequency – Inverse Document Frequency


## Convolutional Neural Network (CNN)

Model Performance on the Test Dataset



Details:

	precision	recall	f1-score	support
0	0.80	0.77	0.79	435
1	0.71	0.75	0.73	327
accuracy			0.76	762
macro avg	0.76	0.76	0.76	762
weighted avg	0.76	0.76	0.76	762



## Model Architecture

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
embedding_6 (Embedding)	(None, 23, 32)	472768
conv1d_6 (Conv1D)	(None, 23, 64)	8256
dropout_3 (Dropout)	(None, 23, 64)	0
max_pooling1d_6 (MaxPooling1D)	(None, 11, 64)	0
conv1d_7 (Conv1D)	(None, 11, 32)	8224
max_pooling1d_7 (MaxPooling1D)	(None, 5, 32)	0
conv1d_8 (Conv1D)	(None, 5, 8)	1032
max_pooling1d_8 (MaxPooling1D)	(None, 2, 8)	0
flatten_5 (Flatten)	(None, 16)	0
dense_9 (Dense)	(None, 256)	4352
dropout_4 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 1)	257

Total params: 494,889

Trainable params: 494,889

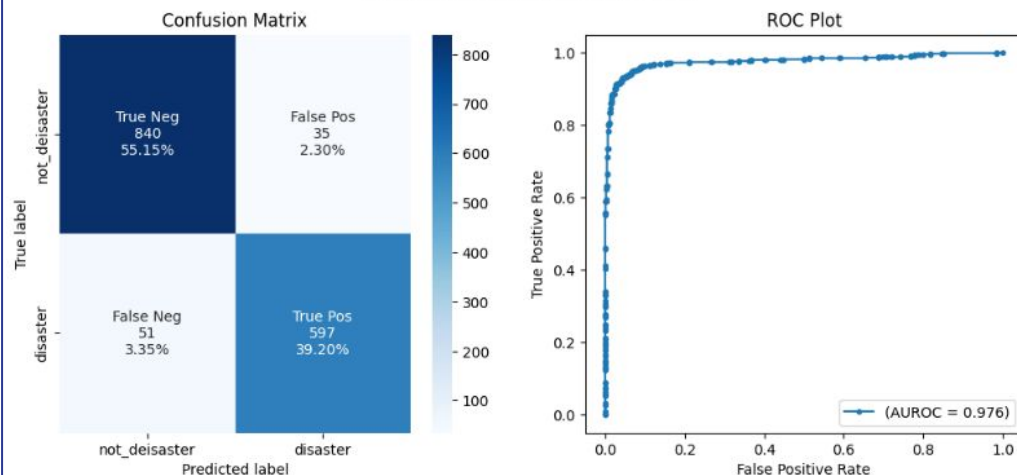
Non-trainable params: 0

# Transfer Learning / Deep Learning

## Bidirectional Encoder Representations from Transformers (BERT)

### Model Architecture - DistilBERT

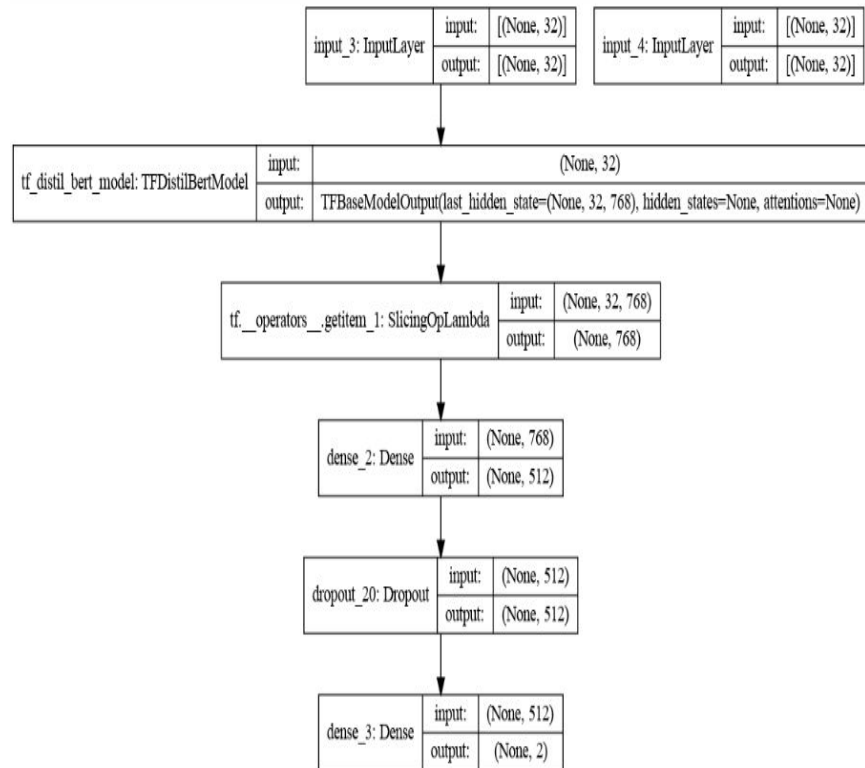
Model Performance on the Test Dataset



Accuracy=0.944  
Precision=0.945  
Recall=0.921  
F1 Score=0.933

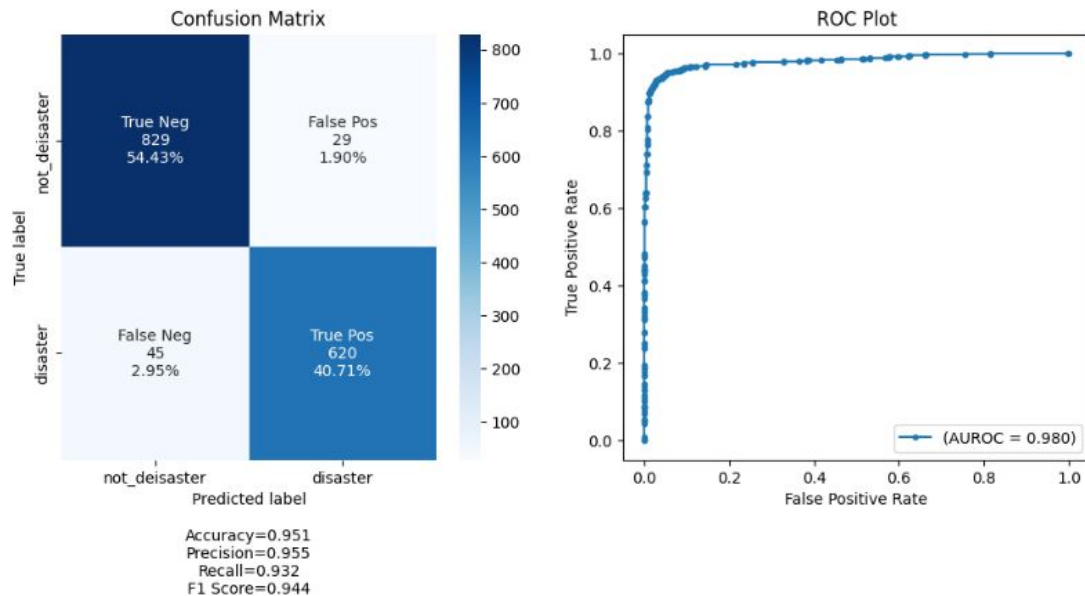
Details:

	precision	recall	f1-score	support
0	0.94	0.96	0.95	875
1	0.94	0.92	0.93	648
accuracy			0.94	1523
macro avg	0.94	0.94	0.94	1523
weighted avg	0.94	0.94	0.94	1523



# Bidirectional Encoder Representations from Transformers (BERT)

Model Performance on the Test Dataset

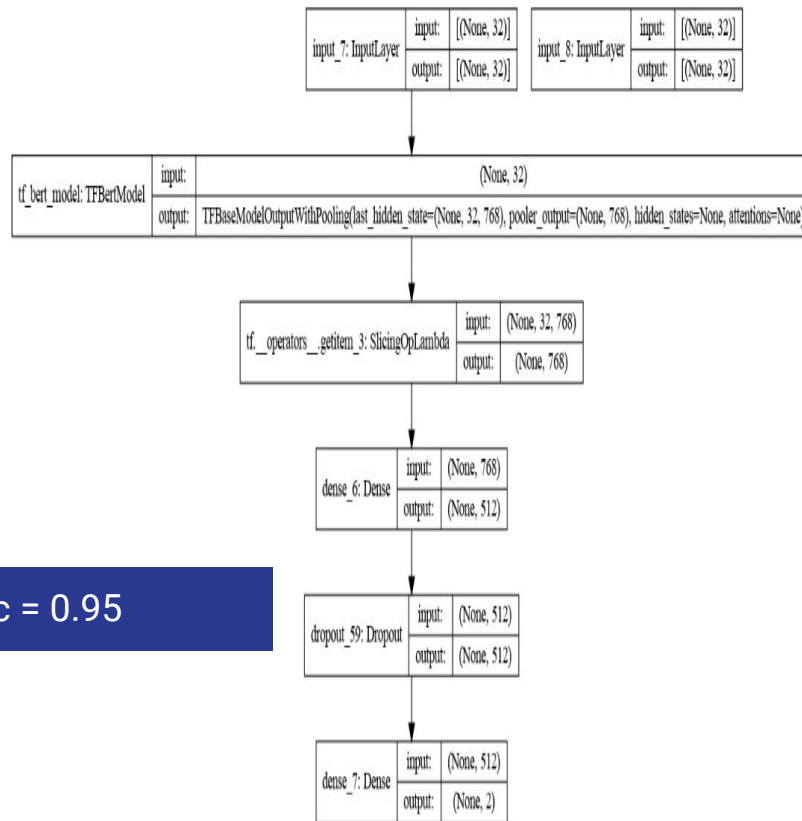


Details:

	precision	recall	f1-score	support
0	0.95	0.97	0.96	858
1	0.96	0.93	0.94	665
accuracy			0.95	1523
macro avg	0.95	0.95	0.95	1523
weighted avg	0.95	0.95	0.95	1523

Best performance Model Acc = 0.95

## Model Architecture - base uncased



## Evaluating the Winning Model on Unlabeled Test Set

	text	Disaster Prediction
0	Just happened a terrible car crash	1
1	Heard about #earthquake is different cities, stay safe everyone.	1
2	there is a forest fire at spot pond, geese are fleeing across the street, I cannot save them all	1
3	Apocalypse lighting. #Spokane #wildfires	1
4	Typhoon Soudelor kills 28 in China and Taiwan	1
5	We're shaking...It's an earthquake	1
6	They'd probably still show more life than Arsenal did yesterday, eh? EH?	0
7	Hey! How are you?	0
8	What a nice hat?	0
9	Fuck off!	0
10	No I don't like cold!	0
11	NOOOOOOOOO! Don't do that!	0
12	No don't tell me that!	0
13	What if?!	0
14	Awesome!	0
15	Birmingham Wholesale Market is ablaze BBC News - Fire breaks out at Birmingham's Wholesale Market	1
16	@sunkxsedharry will you wear shorts for race ablaze ?	0
17	#PreviouslyOnDoyinTv: Toke Makinwa's marriage crisis sets Nigerian Twitter ablaze...	1
18	Check these out: #nsfw	0
19	PSA: I'm splitting my personalities.\n\n?? techies follow @ablaze_co\n?? Burners follow @ablaze	0
20	beware world ablaze sierra leone & guap.	0
21	Burning Man Ablaze! by Turban Diva via @Etsy	0
22	Not a diss song. People will take 1 thing and run with it. Smh it's an eye opener though. He is about 2 set the game ablaze @CyhiThePrynce	0
23	Rape victim dies as she sets herself ablaze: A 16-year-old girl died of burn injuries as she set herself ablaze	1
24	SETTING MYSELF ABLAZE	1
25	@CTVToronto the bins in front of the field by my house wer set ablaze the other day flames went rite up the hydro pole wonder if it was him	1
26	#nowplaying Alfons - Ablaze 2015 on Puls Radio #pulsradio	0
27	'Burning Rahm': Let's hope City Hall builds a giant wooden mayoral effigy 100 feet tall & sets it ablaze. @John_Kass	0
28	@PhilippaEilhart @DhuBlath hurt but her eyes ablaze with insulted anger.	0
29	Accident cleared in #PaTurnpike on PATP EB between PA-18 and Cranberry slow back to #traffic	1

# Conclusion and Performance Evaluation

- Both BertDistill and Bert\_large\_uncased have produced the best performance.
- Bert\_large\_uncased is the winning model with slightly better performance.
- However, training Bert\_large\_uncased with 109,876,994 trainable parameters for 30 epochs required more than 9 hours on 16GB Ram / 4 cores CPU laptop.
- Next we verify the winning model performance on the unlabeled tweets in the test dataset.

- Although Distilled Bert provided slightly less prediction performance, it is faster considerably faster than Bert base model.
- For predicting a streaming tweets in realtime, scalability and latency are a big deal. Depending on the production environment, a trade-off need to be considered between slightly more performance and faster execution time.

